

DENICEK: Computational Substrate for Document-Oriented End-User Programming

Anonymous Author(s)

ABSTRACT

User-centric programming research gave rise to a wide range of compelling programming experiences, including local-first collaborative editing, programming by demonstration, incremental recomputation, schema change control, end-user debugging and concrete programming. Those experiences advance the state of the art of end-user programming, but they are hard to implement on the basis of established programming language and system paradigms.

We contribute Denicek, a system that reduces the complexity of implementing the aforementioned programming experiences. Denicek represents a program as a series of edits that construct and transform a document consisting of data and formulas. Denicek provides three primitive operations on series of edits: application of edits, merging of histories, and conflict resolution. We show how these three operations can be composed to implement a range of programming experiences. We call Denicek a *computational substrate* because it serves as a shared platform for such varied programming experiences.

We discuss the architecture of Denicek, document key design considerations and elaborate the implementation of the programming experiences listed above. To evaluate the proposed architecture, we use Denicek as the basis of a simple end-user data exploration environment. The case study shows that the Denicek computational substrate provides an appealing basis for the design of rich, interactive end-user programming systems.

This is a very promising paper. It's so much simpler and more complete than my stuff. I need to consider some of the different choices you've made. I think I am finally getting what your code execution model is. I hope the comments below are helpful.

KEYWORDS

Programming Systems, Computational Substrate, End-User Programming, Programming by Demonstration, Local-First Software

ACM Reference Format:

Anonymous Author(s). 2025. DENICEK: Computational Substrate for Document-Oriented End-User Programming. In *Proceedings of (Conference acronym 'XX)*. ACM, New York, NY, USA, 17 pages. <https://doi.org/XXXX.XXXX>

1 INTRODUCTION

A computational substrate defines the basic structures with which programs are constructed, how the program state is represented, and how it is changed [21]. The choice of a substrate affects what programming experiences can be readily supported by a system. For example, object-oriented programming has been historically linked to graphical user interfaces [25], while representing programs as lists enabled Lisp to become a “language laboratory” [60].

In principle, any programming experience can be developed on top of any computational substrate. However, a suitable program representation and operations for working with programs provided by the substrate can reduce the complexity of implementing interesting programming experiences. For example, the reflective capabilities of Smalltalk make it easy to provide rich debugging tools [54] that are difficult to provide for C/C++ [26, 27].

Programming Experiences. In this paper, we aim to design a programming substrate that makes it easier to develop a range of rich, interactive programming experiences [43]:

- *Local-First Collaboration.* Multiple users should be able to concurrently modify a single document and automatically merge their changes, preferably without requiring a central server [30, 38].
- *Programming by Demonstration.* Allow users to construct simple programs by demonstrating the steps of the expected behavior using concrete examples [7, 36].
- *Incremental Recomputation.* When a part of a document changes, formulas whose result depends on it are invalidated and, possibly, automatically recomputed [19, 41, 51]. *cite classic incremental computation literature*
- *Schema Change Control.* When the user evolves the structure of the document, affected data and formulas should automatically co-evolve to match the new structure [15, 39].
- *End-User Debugging.* The user should be able to ask provenance questions [6] – why a computation resulted in a particular value and what inputs contributed to the result [33].
- *Concrete Programming.* It should be possible to reuse parts of program logic, or formulas, without introducing abstractions, that is, program against concrete values [13, 14]. *find more references*

Two-Stage Methodology. Our approach is captured by the interior mode [1] of design science research. To *design* Denicek, we identify six *formative examples* – simple programming tasks that manifest one or more of the desired programming experiences (§A). Using those examples, we co-design the Denicek substrate and Webnicek (§3), a simple web-based end-user programming environment based on the substrate. Although Webnicek can be used to complete end-user programming tasks, it is optimized for testing the underlying substrate rather than for usability.

To *evaluate* Denicek, we use the substrate as the basis of Datnicek (§7), an end-user data exploration environment inspired by existing systems [9, 24]. We use the second stage to evaluate suitability of the Denicek substrate for the development of end-user programming systems and report the results of our evaluation (§8).

Substrate. The Denicek substrate brings together two central design ideas. First, it represents programs as document trees consisting

of nodes that can represent data, formulas, evaluated results, as well as static content. Second, Denicek does not store the document tree itself, but instead, it maintains a sequence of edit operations through which the tree was constructed and transformed.

The substrate then provides three primitive operations for working with sequences of edits. First, it can apply a series of edits to reconstruct a document. Second, it can merge two diverging edit histories. Finally, it can detect conflicts when merging histories and, for example, remove conflicting edits from one branch.

The key insight that we present in this paper is that many of the programming experiences can be obtained by leveraging the uniform document representation alongside with a suitable combination of the three primitive operations.

Changing data and formulas is done using the same primitive edit operations to manipulate the document structure. A user-interface may provide a specialized editor, but still trigger the primitive edits behind the scenes. However, interacting with elements in the document, such as entering text in a textbox, also generates a document edit that can be merged or checked for conflicts (§5.3).

As we will see, past edits that demonstrate an operation done to the document can be recorded, allowing programming by demonstration (§5.2). Replaying such recorded edits is implemented using the merging operation, which means that recorded operations continue working even if the document structure later changes. Moreover, structural changes to the document can be done independently and then merged (§5.1). Evaluation of formulas also generates document edits (§5.4). If the evaluated edits conflict with manual edits done later by the user, the evaluated edits are removed, implementing an incremental recomputation (§5.5).

Contributions. The structure and contributions of this paper are:

- We present the Denicek substrate (§4) and provide a detailed description of its document representation, edit operations and the key three operations.
- We explain how a wide range of end-user programming experiences can be implemented using the substrate (§5) in the context of a simple web-based prototype (Webnicek).
- We document important design decisions and alternatives (§6). The analysis shows that the desired functionality requires a careful choice among interconnected design options. **We also identify problems requiring future work.**
- We develop an end-user data exploration system (Datnicek) on top of Denicek (§7) and evaluate the extent to which the Denicek substrate enables such development (§8).

To enable others build on top of Denicek, we share our compact and documented source code at: <https://github.com/d3sprog/denicek>

2 BACKGROUND

What is this about? end-user programming [43] but also programming systems [20], Live [55], Interactive etc. - perhaps for a lack of better term use end-user (also notebooks)

explain computational substrate Substrate defined by [21]

citizen programmers? Maybe do not talk about “programs” because this is more documents with formulas - not that fancy programs. Computational documents?

programming by demonstration [7, 36] [5]
 local first [30, 31]
 provenance/debugging [32–34] [49, 56] [50]
 end-user debugging [29]
 live/incremental recomputation [19, 41, 51].
 visualizing results of execution projection boxes [37]
 livelits (provide editors) [46]
 edit calculus (based on tree navigation though) [47]
 structure editing see [2]
 Schema evolution
 Also on substrate see

Webnicek follows the *naive realism* [8] principle and makes the entire document visible to the user, although parts can be collapsible or hidden using CSS.

Merging - c.f. mywebstrates - we need more than Automerge (cite) which has a bunch of pre-built structures

[16]

The Denicek substrate owes much to systems discussed in §2, especially Subtext, BootstrapLab and Infra [12, 18, 21].

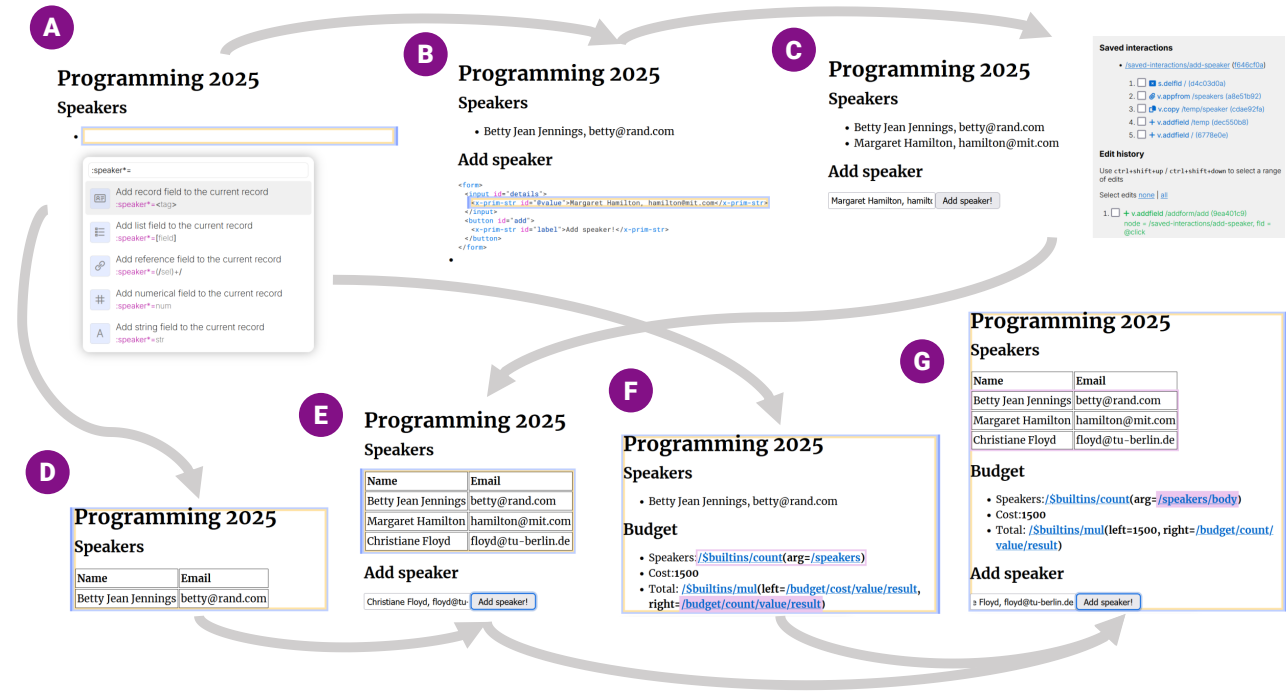


Figure 1: Organizing a conference using Denicek. The Walkthrough shows construction of a user interface for adding speakers (A, B, C); refactoring of the list and merging edits (D, E); and formulas with schema and code co-evolution (F, G).

3 WALKTHROUGH

In this section, we use the web-based prototype Webnicek to illustrate programming experiences enabled by the Denicek substrate (§5). Webnicek is based on a structure editor that supports navigation in the document and issuing of edit commands. We follow two formative examples (§A) related to conference planning.

A Adding a Speaker. The user starts with an empty document, which is represented as a record. They add a field for each heading and a list ``, represented by a field named `speakers`. They use the command toolbox to add the first speaker.

B Creating a User Interface. To simplify adding of further speakers, the user creates a textbox and a button. They enter new speaker details, add a new `` element to the list and copy the speaker details from the textbox into the new element using source view.

C Abstracting the Interaction. After adding the speaker from the textbox, the user opens history view, selects edits that added the speaker and saves those as the `add-speaker` interaction. They attach this as a handler for the `click` event of the button.

Programming by Demonstration. Denicek implements programming by demonstration (§5.2) by letting the user to save and replay past interactions, either directly or by attaching them to a UI event.

D Refactoring Document Structure. Another user starts with the initial version of the document and turns the list into a table. This is done by invoking a series of commands that change tags, wrap elements, copy and transform values.

E Merging Edits. The two document versions can be automatically merged. The refactoring is applied to all existing speakers. New speakers added using the “Add speaker!” button are also automatically added in the new format.

Local-First Collaboration. Denicek’s merging reapplies edits from another branch on top of the current history (§5.1). Merging is asymmetric, but the order does not matter in the above scenario. The same merging operation is used when handling user interaction (§5.3).

F Adding Budget Calculation. A third user adds budget calculation to the initial document. Formulas are represented as special `x=formula` nodes whose arguments are other nodes or references to other nodes or formulas (highlighted on mouse hover).

Incremental Recomputation. Evaluating a formula yields additional edits that augment the document with the result. Those edits are kept at the top of the document history and are removed in case of a conflict (§5.5), providing incremental recomputation.

G Merging Formulas. When the budget calculation is merged with the other edits, references in formulas are automatically updated to point to the new list. Adding new speaker via the “Add speaker!” button invalidates the evaluated result.

Schema Change Control. The substrate understands references in the document and updates them when applying structural edits (§5.6). Evaluation can replace formulas with values, but also augment them to enable end-user debugging via provenance analysis (§5.7).

Selector	Notation	
Parent	..	Refers to a parent of a node
Field	field	Refers to record field of a given name
Index	#index	Refers to list element at a given index
Any	*	Refers to all children of a list node





Kind	arguments	
	List <i>tag, index₁, child₁, ..., index_n, child_n</i>	Ordered list of nodes, addressable by <i>index</i> . Renders as <tag> with children.
	Record <i>tag, field₁, child₁, ..., field_n, child_n</i>	Record with children addressable by <i>field</i> . Renders as <tag> with children.
	Reference <i>selectors</i>	Reference to another document location. Displays the selectors as a link.
	Primitive <i>string or number</i>	Numerical or textual primitive value. Renders as an HTML text node.

Figure 2: Structure of selectors and document nodes

4 THE DENICEK SUBSTRATE

Denicek represents programs as sequences of edits that construct and transform a computational document. In this section, we describe the structure of documents and edits, as well as the operations that form the backbone of the system and are used to implement the end-user programming experiences discussed in §5.

4.1 Selectors, Documents and Edits

A computational document is a tree, consisting of four kinds of nodes (Fig. 2). Non-leaf nodes are records or lists. Both of these are named with a *tag* as in HTML. Record fields have **unique?** names. Lists contain elements of the same type and assign a unique index to each element **It seems that these indexes are meant to be unique IDs, in which case it would be clearer to call them that.** Leaf nodes can be primitive values (strings and numbers) or references to a location in the document tree. References can be relative or absolute. Both kinds can appear in a document (reference nodes); absolute references are used in edits (to denote a target node). **Why are relative references needed? I only use absolute references.**

A reference is represented as a sequence of selectors (Fig. 2). The document model assumes that lists are homogeneous and records heterogeneous, and so the Any selector makes it possible to refer to all children of a list, but there is no way to refer to all children of a record. Note that Denicek does not use implicit numerical indices for lists. The index of a new item has to be supplied explicitly. The reasoning behind this design choice is discussed in §6.

Document Edits. The supported document edits and their behavior are listed in Fig. 3. All edits require *target* to which they are applied. Target is an absolute reference not containing the Parent selector. It can contain the Any selector, in which case the edit is applied to multiple nodes simultaneously. Most edits can only be applied to target node(s) of a specified kind. As discussed in §6, both fields of a record and list elements are ordered and edits that add a new item take the index or field name of a previous node.

The edits allow any transformation of a document through a series of steps whose effect can be tracked by the substrate. As we saw, Denicek updates references when document structure changes.












Edit	arguments	Target
	Add <i>target, field, after, node</i>	Record
	Add <i>node</i> as a <i>field</i> to the specified record <i>after</i> a given field.	
	Append <i>target, index, after, node</i>	List
	Append <i>node</i> to the end of the specified list <i>after</i> a given field.	
	Reorder <i>target, permutation</i>	List
	Reorder items of a specified list according to a <i>permutation</i> .	
	DeleteItem <i>target, index</i>	List
	Delete the item at a given <i>index</i> of a specified list.	
	UpdateTag <i>target, tag</i>	List or Record
	Change the tag of a specified list or record from to a new <i>tag</i> .	
	PrimitiveEdit <i>target, transform</i>	Primitive
	Apply primitive <i>transform</i> to the specified primitive.	
	RenameField <i>target, old field, new field</i>	Record
	Rename the field of a specified record from <i>old</i> to <i>new</i> .	
	DeleteField <i>target, field</i>	Record
	Delete the field <i>field</i> of a specified record.	
	WrapRecord <i>target, tag, field</i>	Any
	Wrap the specified node as a <i>field</i> of a new record with <i>tag</i> .	
	WrapList <i>target, tag, index</i>	Any
	Wrap the specified node as a sole element of a new list with <i>tag</i> .	
	Copy <i>target, selectors</i>	Any
	Copy nodes(s) from <i>selectors</i> , replacing the specified target(s).	

Figure 3: Summary of document edit types in Denicek

Fig. 3 distinguishes between edits that keep references in a document unchanged (above) and edits that affect references (below).

Renaming a field or wrapping a node converts any references into that location (§4.2). When deleting a field to which there is a reference in the document, Denicek rejects the edit. A copy edit can also be rejected, because it is ambiguous whether references referring to the original location should continue referring to the source or should be changed to point to the target of the copying (a reference cannot refer to two locations). **why not have distinct Copy and Move operations to resolve this?**

Automatic Reference Update. There are two situations in which automatic update of references is undesirable. If an edit is applied to a singular element of a list (reference contains the Index selector), references that refer into the list should be unchanged. Although such edits violate the invariant that collections are homogeneous, they typically do so only temporarily, for example during construction of a new list item (checking such edits is discussed in §6.1). **Wouldn't this leave dangling references? Wouldn't it be better to evolve the references as these edits occur? Insertion into a list can default the contents based on the type to avoid such a transient type error.**

Documents can also contain values that can be of multiple different kinds (i.e., a union type). **This is surprising since there has been no prior mention of union types, and lists are homogeneous.** In such cases, references should not be updated when the kind of the value changes. For example, a formula may be either unevaluated or evaluated. As discussed in §5.4, evaluation involves wrapping the formula node, but this should not affect references to the formula. To support those cases, edits that normally affect selectors

(Fig. 3) have a *reference behavior* field that can be set to disable automatic reference updating. This annotation is required when the target contains the *Index* selector. I think it would be desirable to avoid all these complications. References should be automatically maintained so they “just work”. I believe that evaluation should not change the type/shape of the document. The space for evaluation results should be preallocated in some sense, either as extra structure in a formula, or as metadata attached to the tree. That said, I’ve tried various alternative approaches to this in Subtext without ever settling on a preferred solution.

Edits are “static”. An important principle is that the effect an edit has on references inside the document does not depend on the current value of the document, only its type. This makes it possible to define merging solely in terms of edits, without reference to current document state. (The effect Copy has on the document structure depends on the existing document structure, but not on its value).

This design choice makes it impossible to encode computational logic directly in the edits (e.g., through conditional edits). As we will see in §5.3, such logic can be provided as an additional mechanism on top of the underlying Denicek substrate.

4.2 Primitive Operations

Denicek provides three primitive operations. A sequence of edits can be applied to a document, two sequences of edits can be merged and also checked for conflicts. Denicek identifies edit histories by a (git-like) hash, computed from the hash of the parent and the current edit. The hash is used to identify a common shared part of the history during conflict resolution and merging. *what exactly does parent mean here? History prefix?*

Applying Edits. When applying an edit, Denicek locates the target node and transforms it according to the edit. If the edit can affect references in the document (Fig 3, below), Denicek updates matching references in the document according to the rules shown in Fig. 4 provided that the reference behavior of the edit is not set to disable the updating (this is required when the target reference contains the *Index* selector). Reference update behavior for WrapRecord and WrapList differs in that the updated references as a result of WrapList use the *All* selector. Although WrapList specifies the index to be used for the newly created list item, we assume that the operation introduces a homogeneous list and the new reference should point to all eventual list items.

Also note that affected references in the document may be more specific than the edit target. For example, if we rename *old* to *new* at */foo/**, a reference */foo/3/old* will become */foo/3/new*. (A reference in the document cannot be more general, because such edit would have to contain the *Index* selector and this would require setting reference behavior to not trigger reference update.) Finally, if the document contains a reference that would be invalidated by the Copy or DeleteField edit, the edit is rejected.

Merging Edit Histories. Merging edit histories is used when two users edit document independently, but also when replaying edits in programming by demonstration. The operation works on two edit histories, E, E_1 and E, E_2 , that have a shared prefix E . The merging operation is akin to git rebase. It turns edits E_2 into edits E'_2 that

RenameField	<i>target, old field, new field</i>	– Replace Field for matching references. <i>/target/old_field/nested</i> \Rightarrow <i>/target/new_field/nested</i>
WrapRecord	<i>target, tag, field</i>	– Insert extra Field selector after matching prefix. <i>/target/nested</i> \Rightarrow <i>/target/field/nested</i>
WrapList	<i>target, index, tag</i>	– Insert extra All selector after matching prefix. <i>/target/nested</i> \Rightarrow <i>/target/*/nested</i>

Figure 4: How document edits transform references

can be reapplied on top of the other edit history, resulting in a new history E, E_1, E'_2 . Note that the operation is not symmetric. If there are conflicts among the edits in E_1 and E_2 , the result of E, E_1, E'_2 will differ from the result of E, E_2, E'_1 . We return to this problem in the next section when discussing conflict detection.

It might help to define a function signature for merging to make this clearer: $\mathcal{M}(P, L, R) = M$ or $\mathcal{M}_P(L, R) = M$. Merging is often thought of as symmetric, perhaps choose a name that is clearly asymmetric, like *rebase*, *translate*, or *project*.

The key operation that enables such reconciliation takes two individual edits that occurred independently, e_1 and e_2 , and produces a sequence of edits e'_2, e''_2, \dots that can be applied after e_1 and have the effect of e_2 , modified to respect the effects of e_1 . There are two aspects of such reconciliation:

- (1) *Apply to Newly Added.* If e_2 is adding new nodes to the document, but e_1 modified the document through a selector that would also affect the new nodes added by e_2 , we need to apply the transformation represented by e_1 to the nodes newly added by e_2 . This is done by generating an additional edit, to be applied after e_2 , that is based on e_1 but targets only the newly added nodes (more details can be found in §B.1).
- (2) *Transform Matching References.* If e_2 targets a node that is inside a node whose structure is changed by e_1 , the target reference in e_2 is updated in a way that corresponds to the new structure. This is done using the rules shown in Fig. 4, that apply when transforming references inside a document, although we also support the case when e_1 is Copy (see §B.2 for details).

I believe that in general a “backwards” form of merge is also needed, as in my *retract* and the *OT Exclusion Transform*. A regular merge takes forking edits and composes them. The backwards form converts composed edits into forking edits (and can be combined with a forward merge to recompose in the opposite order). This is needed for cherry picking and reverts, however you don’t seem to need that for this paper.

Related question: how do you merge an `Append(target, index, after, node)` with `DeleteItem(target, after)`? A famous counterexample in OT was that two insertions adjacent to a deletion could flip their location if the edits were reordered. I’ll dig up the details of this counterexample for you. Some systems handle this by leaving behind a “tombstone” on delete to remember the position of the deleted item. I have a novel solution but I’m probably going to go back to tombstones anyway because they are simpler and handy in the diff UI.

To illustrate merging, consider a case where we created a list of work items */todo*. In one branch, we add an additional item to the list (e_2). In another branch, we wrap the list in an extra `<div>`

StructureEffect <i>target</i> – Affects fields or structure of the target node. RenameField, DeleteField, WrapRecord, WrapList, Copy
ValueEffect <i>target</i> – Transforms value, modifies list or adds an additional field. Add, Append, Reorder, DeleteItem, PrimitiveEdit
TagEffect <i>target</i> – Modifies the tag of the target node. UpdateTag

Figure 5: Effects of individual edit operations

element (e_1) and add a checkbox to each work item (e'_1):

$e_2 = \text{Append}(/todo, \#1, \#0, \text{Do some work})$

$e_1 = \text{WrapRecord}(/todo, <div>, \text{items})$

$e'_1 = \text{Add}(/todo/items/*, \text{done}, \text{nil}, <\text{input type}=\text{"checkbox"}>)$

If we want to append the edit e_2 after edits e_1, e'_1 , we need to update its target to reflect the additional wrapping (2) and we need to create an additional Add operation that will add the checkbox to the new item (1). The result is two edits e'_2, e''_2 :

$e'_2 = \text{Append}(/todo/items, \#1, \#0, \text{Do some work})$

$e''_2 = \text{Add}(/todo/items/\#1, \text{done}, \text{nil}, <\text{input type}=\text{"checkbox"}>)$

If we performed the merge operation in the other order, we would simply append e_1, e'_1 after e_2 . Although this is not the case in general, the result will be the same regardless of the order here.

Conflict Resolution. When merging two sequences of edits, E, E_1 and E, E_2 , it is sometimes possible to construct edit histories E'_1 and E'_2 such that the results of applying E, E_1, E'_2 and E, E_2, E'_1 are the same. However, this is not always the case. Two edits can conflict if they both modify the same value, if they transform the structure of a node in incompatible ways or if one modifies a node nested inside a node replaced or deleted by the other. Detected conflicts can be reported to the user, but we also use conflict detection to implement incremental evaluation (§5.5). Our treatment of conflicts necessarily differs in this respect from replication algorithms like Operational Transformation [?] and Convergent Replicated DataTypes [?] where conflicts are arbitrarily but deterministically resolved.

The Denicek substrate implements a conflict detection mechanism inspired by effect systems [40]. The mechanism is simple and tractable, but over-approximates conflicts, i.e. it may report a conflict even if two edits can be merged successfully (see §6.1). An effects describe how an edit affects the document structure and we distinguish between three types of effects as shown in Fig. 5.

We say that a set of effects F_1 conflicts with another set of effects F_2 if there are effects $f_1 \in F_1$ and $f_2 \in F_2$ such that they are of the same kind and the target of f_1 is a prefix of the target of f_2 or vice versa (allowing specific Index to match against All in any direction).

Given two edit histories E, E_1 and E, E_2 , Denicek can use conflict detection to remove all conflicting edits from E_2 and produce a sequence of remaining edits e'_2, e''_2, \dots that can be added after E, E_1 and do not conflict with edits in E_1 .

To do this, we iterate over edits e_2 from E_2 and check if the dependencies of e_2 conflict with effects of (1) any of the effect e_1 from E_1 or (2) effects of any of the previously removed edits. Here, the dependencies of e_2 include its target, but also source of Copy and additional dependencies recorded explicitly as discussed in §5.5. If a conflict is detected, the edit e_2 is removed and its effect

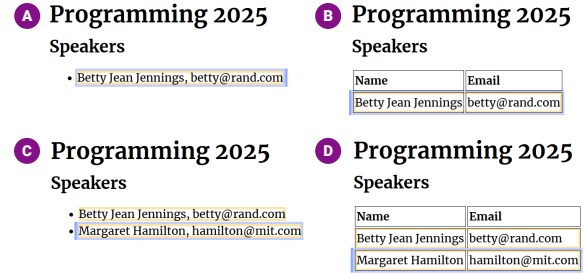


Figure 6: Merging of two independently done sequences of edits. Two ways of merging B and C result in the same D.

is recorded, so that we remove any subsequent edits that would depend on the removed edit.

The effect-style conflict detection is a nice idea! Removing conflicts before merging is also an interesting idea I hadn't considered. That might allow me to simplify my transformation rules.

5 PROGRAMMING EXPERIENCES IMPLEMENTATION

The key claim of this paper is that the Denicek computational substrate supports a range of experiences that make programming more concrete, collaborative, and interactive. In particular, Denicek's ability to merge edits is a unifying mechanism that provides essential features across the range. We demonstrate this claim with the Webnicek web-based programming system, which uses Denicek to support:

- (1) Local-first collaboration (§5.1)
- (2) Programming By Demonstration (§5.2§5.3)
- (3) Incremental recomputation (§5.4, §5.5)
- (4) Schema change control (§5.6)
- (5) End-user debugging via provenance tracking (§5.7)
- (6) Concrete programming via managed copy & paste (§5.8).

We describe the programming experience in isolation in this section. The next section provides a more comprehensive evaluation through a case study that combines multiple of them together.

5.1 Local-First Collaboration

The Denicek representation enables local-first collaboration [30], as illustrated previously in §3 (E). If a document is edited by multiple users, they can each make edits to their local copy and eventually merge the variants using the operation to merge edit histories.

Merging of histories behaves akin to git rebase in that it keeps a linear history. Synchronization in a distributed system thus requires first reapplying local edits on top of the remote history, before updating the remote history. Denicek thus implements the *convergence model* of document variants [15], i.e., the user cannot, for example, maintain their own local document structure and import new data from another variant (an alternative discussed in §6).

Recall that merging of edit histories is not symmetric. Document D in Fig. 6 can be obtained either by appending C' (produced by the edit reconciliation operation) on top of A, B or by appending B' on

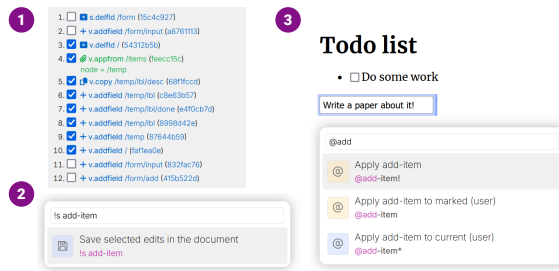


Figure 7: Programming by demonstration is implemented by selecting edits from the document history (1), saving them in the document (2) and replaying them (3).

top of *A*, *C*. According to our effect analysis, the two operations are conflicting. Although *B* primarily affects the document structure, it also adds a new field to the record (email), which is a ValueEffect, conflicting with ValueEffect of *C*. In this scenario, the conflict can be ignored and the resulting document is the same in both cases.

However, the resulting two histories will differ. In the first case, the Append edit that adds a new node is supplemented by further focused edits that transform the added node from a list item to a new table row. In the second case, the structural transformations are automatically applied to all rows. As discussed in §4.2, conflicts during merging can be resolved either by removing conflicting edits or by letting the later edits overwrite the former ones.

To guarantee that replicas converge it is necessary to have a global order on merges. OT and CRDT algorithms use timestamps combined with a unique ID for each client to establish this global order without a centralized coordinator. We can do the same. This means you never know if your changes will dominate or whether they will be eventually overridden but people seem to be willing to live with it. Personally, I'm a skeptical: I expect that like with NoSQL people will eventually realize this sucks and go back to centralized systems.

5.2 Programming By Demonstration

In programming by demonstration [7], the user demonstrates a task to the system and the system then repeats it, directly or in a generalized way. To use direct repetition with Denicek (Fig. 7), the user can select edits from the edit history, name them and replay them. In case of general-purpose document editing in Webnicek, this requires certain forethought, but as illustrated in §7, the mechanism is very effective in a domain such as data wrangling [24].

There are two notable aspects of our implementation of programming by demonstration in Webnicek. First, Webnicek records edits in the document itself (by representing individual edits as nodes and storing them in list inside a /saved-interactions field). This means that no other implementation mechanism outside of the system is needed and also that the stored edits can be modified by the user or tools working with the document.

Second, to replay edits, Webnicek does not append the recorded edits on top of the current history. When saving edits, it stores the hash of the history at the time of saving. When replaying, it appends the recorded edits to the top of the original history (at the time of




Figure 8: Using programming by demonstration to define a UI. The “Add” button (1) replays edits; the “Remove completed” button (2) modifies target and specifies a condition.

saving) and merges this new sequence of edits with the current history. This pushes the recorded edits through all subsequent edits made by the user. The result can be seen in Fig. 1 (E), where a newly added speaker is transformed from a list item to a table row.

Webnicek also accounts for the case where edits recorded in the document are themselves transformed (when they are reconciled with other edits during merging). In this case, Webnicek updates the recorded edits (an alternative approach is discussed in §6).

Simply replaying recorded edits provides limited programming capabilities. We discuss in the next section how recorded edits can themselves be edited by the user to generalize and conditionally their effects. More general and conventional programming constructs are discussed in §6.

5.3 Interactive User Interfaces

Programming by demonstration can be used to define interactive elements in the document. In a simple scenario, illustrated in Fig. 8 (1), the click event handler is set to a reference to a sequence of edits recorded in the document. Clicking the button executes the edits using the mechanism discussed in §5.2, i.e., Webnicek appends the edits to a history at the time when the edits were recorded and merges the edits with the current history.

The use of merging when replaying recorded edits is crucial in both the Todo App example and the Conference List example (see §A). In both cases, it makes it possible to define a user interface for adding items (new Todo items, new speakers) and later change the document structure (refactor speaker list to a table) or add functionality (formula to evaluate whether a Todo item has been completed) without having to recreate the user interface for adding items. The use of merging ensures that new items are added in a correct format or with the additional functionality.

Denicek can be used to generalize the interactions recorded through programming by demonstration. Our prototype illustrates this option in a limited way. As shown in Fig. 8 (2), a button to remove all completed Todo items can be created by generalising the remove-item interaction, which removes the list item at the index 0. In addition to the recorded interaction, we manually specify (in the source view) that the edits should be applied to all elements selected by the /items/* selector, instead of the original /items/0 selector (prefix) and that the edits should only be applied to elements for which the formula (which tests if the checkbox is checked) specified by a relative selector ./condition/comp/result evaluates to true.

The above paragraph was surprising to me and I'm not sure I fully understand it even after staring at the figure for a while. Edits are recorded as an AST using x-tags to hide them in the UI? That became clear only after reading the next section. Our audience might be a lot happier if we used a textual grammar. Earlier it was said that edits can't be conditional but then a conditional predicate appears here. This is a special feature of button nodes? Or event handler nodes? This starts to resemble Angular and similar template languages (which I find hard to read).

✎ **Generalization Heuristic.** Specifying generalization manually is cumbersome. Programming by demonstration systems typically implement heuristic for generalization [44] infers and suggests such generalizations. If integrated into a system based on Webnicks, such heuristic could automatically construct a formula based on positive and negative examples [35] (selected and deselected items).

5.4 Formula Language and Evaluation

Denicek documents can contain formulas inspired by the spreadsheet paradigm [45]. Formulas can specify richer computations than what can be expressed using document edits. Formulas do not transform the document and their results are transient.

Formula evaluation leverages Denicek’s ability to merge edit histories. As illustrated in Fig. 9, formulas are represented as document nodes with a special tag (`<x-formula>`). They are recognized by a formula evaluator and rendered in a special way (❶ and ❹), but they are created using ordinary edits and the Denicek substrate treats them as standard nodes.

To evaluate formulas, the formula evaluator generates edits that turn the `<x-formula>` record into `<x-evaluated>` (2 and 3), which keeps the previous formula state in the `formula` field and the evaluation result in the `result` field. (Keeping the previous state of the formula is not necessary, but it enables provenance analysis as discussed in §5.7.) The way edits generated by evaluation are merged with the document is discussed in the next section.

As I've said before, I'd prefer to store the result of the formula inside the <x-formula> record instead of wrapping it on evaluation. This avoids changing the shape of the document on evaluation, which I recall added some complexity earlier.

The Counter App example shown in Fig. 9 illustrates the interaction between formulas and programming by demonstration. To implement a counter, the Increment and Decrement buttons wrap the current counter value in a formula that adds or subtracts 1. I'd like to see the code for the increment button.

OK maybe I'm finally getting it. You aren't mutating the value of the counter, you're wrapping the value in an increment formula. I never got that in our previous discussions. Presumably that evaluation is now frozen so it serves as an execution trace. This is a lot different from my approach where a formula stays unchanged except for the results. One advantage is that fields contain a record of their computed mutations whereas in my approach that is recorded in the past state of the doc whence the mutation occurred.

x² Formula Language. Webnicek exposes the underlying representation of formulas to the user, but the same representation can be edited through a user-friendly mechanism such as a block-based editor

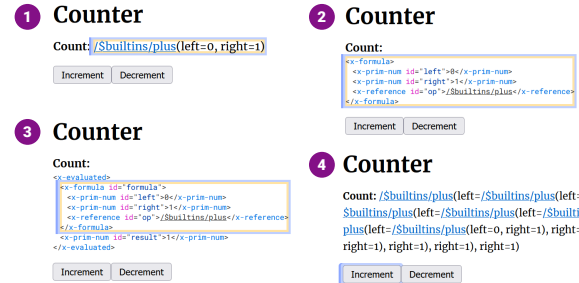


Figure 9: Increment wraps the existing count in a formula that adds 1 to the previous value (1), (2). Evaluation produces the count (3), which is invalidated on subsequent clicks (4).




Figure 10: Evaluated edits (A, B) are kept at the top of the history. When they conflict with ordinary edits (4), they are removed (A).

[22] or a calculation view [57]. The key point is that Denicek’s tree structure makes it easy to embed formulas in document in a uniform way, edit them and merge them with other changes. **Some readers may disagree that it is easy: AST editing is controversial.**

5.5 Incremental Recomputation

As illustrated in Fig. 11, Webnicek supports incremental recomputation. In the example, the cost of speaker travel depends on the number of speakers, but the cost of refreshments depends only on two constants. The edit that adds a speaker invalidates only the former computation.

The evaluation mechanism is illustrated in Fig. 10. **This figure is out of order.** When the formulas are evaluated, Webnicek generates *evaluated edits* that are appended to the top of the history. **What exactly are these evaluated edits?** When subsequent edits are made, they are appended after all non-evaluated edits and the evaluated edits are pushed through the newly added edits. If the edits conflict (according to conflict detection discussed in §4.2), affected evaluated edits and edits that depend on them are dropped.

Live Programming. The Webnicek prototype does not automatically evaluate formulas. This makes it easier to understand the evaluation mechanism, but a realistic system based on the substrate could automatically evaluate formulas to provide a live programming experience [52, 55] and use incremental recomputation for performance reasons.

I'm confused. Changing the value of a location conflicts with an edit that evaluates a formula containing a reference to that location? And these evaluation edits are second-class edits that get retroactively dropped to "unevaluate" the formula? I'd prefer what you suggest above: formulas are seemingly globally re-evaluated after every change, with incremental computation a hidden perf optimization.

An important property of incremental computation systems is avoiding redundant execution. Usually this involves a topological sort of the computation dependency graph. We may need to have

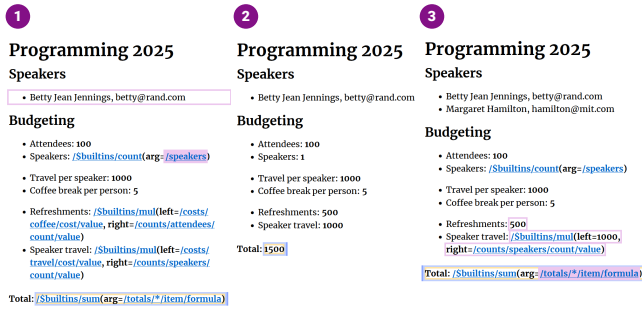


Figure 11: Budget calculation based on the number of speakers (1) and the result (2). When speaker is added (3), only the results of affected formulas are invalidated.

a story about that to avoid hate from reviewer 2. Personally, I think incremental recomputation is about performance, not the programming experience, so it isn't that important for this paper.

5.6 Schema Change Control

Document structure often needs to evolve [4], as illustrated by our Conference List example where a list is transformed into a table. When this happens, data and code that depend on the structure of the document need to evolve correspondingly. The problem is well-known in database systems [53] and has recently been studied in the context of programming systems [15].

Although Denicek does not explicitly track document structure (schema or type), all documents have an implicit structure and some edits transform this structure. As discussed in §4.2, Denicek automatically updates reference nodes in the document when edits modify the document structure. This enables a form of schema code co-evolution [15]. Formulas embedded in Denicek documents use reference nodes to refer to both data sources (in the document) and the results of other computations. Consequently, if the document structure changes, the formulas are automatically updated.

Consider the example in Fig. 12. The original list `` is turned into `<tbody>` using `UpdateTag` and wrapped inside `<table>` with a field `body` using `WrapRecord`. For the latter, Denicek updates the reference accordingly turning the original `/speaker` reference in the formula into `/speaker/body`.

As noted earlier, evaluated edits transform the formula structure (wrapping it in the `<x-evaluated>` record). However, they are marked as non-structural and so references to formulas are not transformed when the formula is evaluated (otherwise, references would be updated to point to the original unevaluated formulas).

5.7 End-User Debugging

The most common kind of end-user programming question is determining whether a value they observe is right or wrong [29]. One way to help users answer the question is to provide an explanation how a value was obtained [33]. Webnicek provides a basic mechanism that highlights document nodes that contributed to a specific computed result, illustrated in Figure 13.

The implementation leverages the fact that evaluation keeps the original formula, as discussed in §5.4. When a formula is evaluated, the final `<x-evaluated>` document node contains the result, but also




Figure 12: When an edit changes the document structure, references in formulas are updated accordingly.




Figure 13: Using provenance tracking to highlight document parts that contributed to the calculation of refreshments costs (1), travel costs (2) and all costs (3).

sub-tree that represents a full evaluation trace [49]. We analyse the trace, collect all reference nodes in the trace and highlight all nodes referred to in the computation. Denicek makes this easy as we only need to collect reference nodes nested in a formula node.



Explanations and Linked Visualizations. In Webnicek, we implement provenance analysis to show inputs involved in computation, but information from the execution trace collected by the Denicek substrate can also be used to provide a detailed explanation [49] or to automatically generate linked visualizations [50].

5.8 Concrete Programming

Abstraction is an essential feature of programming, but it has a high cognitive cost [3]. Programming by demonstration (§5.2) offers one way of reducing the cost. Another way of making programming more concrete [11, 58] is to support copying of functionality as in prototype-based object-oriented programming [59].

Webnicek supports a functionality akin to managed copy & paste [13, 14] for formulas. Rather than introducing abstractions (functions), users can copy and modify formulas to reuse them. However, when the user discovers an error in the original formula, Webnicek lets them use the merging mechanism to correct the error in the original formula and all its copies.

The mechanism is shown in Figure 14. The user copies an incorrect formula using the `Copy` edit and then modifies it to use a different data source. They then navigate back in history to the point before the copying and create a temporary fork of the document. In the fork, they use `RenameField` to switch the argument of the division. When they merge the temporary fork into the original document, the *Apply to Newly Added* logic of the merge operation (§4.2) duplicates the `RenameField` edits and applies them also to the

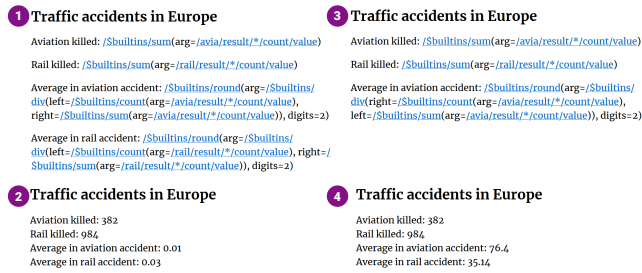


Figure 14: Correcting an incorrect formula. The user uses Copy to create the second formula (1). They notice an error (2), go back in history to switch left and right argument (3), merge the change and re-evaluate both formulas (4).

copied formula. In other words, the Copy edit of the Denicek substrate, alongside with the fact that formulas are ordinary document nodes, provide keys component for a straightforward implementation of the managed copy & paste functionality.

Nice and simple! My approach tried to be a lot fancier but ended up being too heavyweight.



Linked Editing. Webnicek currently requires users to explicitly manipulate history to correct error across multiple code clones. Research on managing duplicated code resulted in multiple tools [10, 61] with dedicated user interface to manage clones. The Denicek substrate provides the underlying mechanism that could be used to implement a more user-friendly interface inspired by those systems.

6 DESIGN CONSIDERATIONS

The design of the Denicek substrate is the result of an iterative process in which we repeatedly adapted the Denicek design until we obtained a satisfactory solution for the six formative examples (§A) in the Webnicek prototype. In this section, we document the design challenges, many of which have until now been personal knowledge of researchers working on related systems [12, 18, 21, 46].

Two guiding principles for Denicek have been *composability* and *uniformity* [20]. The design should cover a large number of cases using a small number of concepts. This is apparent in the design of the document structure (a node can be data, code or rich text) as well as in the role of edits (an edit can be a value or structure change, the result of user interaction or the result of evaluation).

The inconvenience associated with composability and uniformity [20] is not a problem for Denicek. We view Denicek as an underlying structure on top of which more convenient end-user programming systems can be built. The uniform design offers a degree of open-endedness, making it possible to use the basic structures of Denicek in new ways, not anticipated in this paper.

Implicit vs. Explicit Structure. A key design choice for a substrate like Denicek is whether to track the document structure explicitly or make it implicit. In Denicek, we assume that elements of a list node have the same structure (i.e. records should have the same fields, so that they can be targetted using the All selector), but we do not enforce this. Often, the property is temporarily violated when adding a list item, but then restored once the item is created.

Keeping the structure (type) explicit has theoretical appeal and it simplifies some aspects of the implementation. In particular, edits to document structure can be clearly separated from edits of document values. The same separation has disadvantages in that it complicates the edit language and it requires novel user interface. (In Webnicek, we can select all list items and edit them at once using the same mechanism as when editing individual list items. With explicit structure, changing a list structure has to be either a different kind of edit or possibly an edit of a virtual “prototype” element.)

Merging and Document State. One specific design choice in Denicek is that the merge operation operates solely on two sequences of edits, but it does not need access to the current document state (or full edit histories). The reconciliation of edits (§4.2) is defined for two individual edits e_1 and e_2 .

This limits the language of edits in a number of ways. The first consequence is that we cannot, in general, support conditional edits (edits that are applied only when a certain condition holds for the current document state). This is the case because the merging operation would not be able to determine whether the edit is applied (and whether its effects should take place). One consequence of this choice is that generalization of recorded interactions (§5.3) has to be done on a meta-level (by specifying how to invoke edits, rather than in the edits themselves). The second consequence of this design choice is that our language of selectors has to be abstract. We can target all list elements using the All selector, but cannot convert between this representation and a representation that explicitly lists indices of all items currently in the list.

Explicit List Indices. Denicek does not use implicit numerical indices for lists. When adding an item, the user (or the system built on top of Denicek) has to explicitly provide an index. This design is motivated by a typical programming by demonstration scenario where the user adds a new list item and then modifies it (adding a speaker or a Todo item). When using explicit indices, the edits following Append can easily refer to the newly added item. (To avoid overwriting existing items when replaying recorded edits, Webnicek replaces indices of newly added items in recorded interactions in §5.2). *I don't think that was explained back there.*

Using implicit numerical indices is possible, but it requires treating all list-related edits as operations that transform document selectors. This means adding multiple new rules to Figure 4, which specifies how edits transform references. Merging then has to modify indices when edits add, remove or reorder list items. Somewhat, counterintuitively, list order then becomes a part of document structure rather than a property of the list value.

I'm still unclear whether your list indices are ordinal numbers like in an array or are unique IDs. Your usage of implicit/explicit isn't obvious to me. The above paragraph seems to imply they are IDs. Note that automatically generating unique IDs in code execution can lead to nondeterminism. If you are using IDs, do they affect equality of lists?

The problem could also be avoided by not letting the user modify a list item after adding it. (We can require new items to be constructed outside of the list and then added through a single edit that combines the logic of Append and Copy.) This is easier to implement, but it requires unintuitive user interaction and it also

does not address all use cases for modifying list item via index. **I'm still confused, but this sounds like it really isn't a solution afterall.**

Ordering List Items and Record Fields. When list indices are not numerical and supplied explicitly, we need another mechanism for tracking order of list items. Denicek uses a data structure inspired by Mergeable Replicated Data Types [23] and requires specifying the index of a preceding list item when appending to a list. (If two edits append to the same list independently with the same preceding item, the order is non-deterministic.) **Why is this needed? I let the order of operations determine the list order.**

Note that Denicek uses the same mechanism for record fields. Keeping record fields ordered is desirable because Denicek documents map directly to HTML documents and so the order is visible to a user. Moreover, different ways of merging can result in different order of adding fields to a record. Making order explicit ensures those result in equivalent documents.

Denicek still makes a distinction between records and lists, even if the two are technically extremely similar. The difference is conceptual. A list is expected to contain items of the same structure (that can be addressed using the All selector), whereas record fields are expected to be of different types. If we unified lists and records, All would have to be applicable to records too and inferring and checking document structure would be challenging.

two-way merging. divergence model of challenges paper - requires inverse operation for moving edits backwards

keeping git tree. save graph of edits rather than list (would make replaying of saved edits easier because the histhash never changes)
=> PBD - updates the recorded edits this would not be needed if the history was stored as a graph, akin to ordinary git merging rather than rebasing).

ISSUES * if we merge a thing with saved-interactions with something, hashes will change!

Fancy selectors. * many things with <tag> selectors currently do not work (e.g. 'matches' for highlighting) because if we collect path of a current node, we collect indices and get /some/2/another - and cannot tell if this matches /some//another - we'd have to collect more detailed path info!

* selectors - all for structure / index for data (but it is useful to allow others...) (multiselect also bad for checks!)

Dependency tracking. NOTES type Edit = Kind : EditKind Dependencies : Selectors list - only needed for evaluated edits

Maybe have 'enabled' for edits afterall? (we can merge with conflicts and disable some edits, but keep them in history for info)

TODO - things to work on * "represent" edits somewhere in document as "library of functions" and then call those from buttons (rather than embedding them directly) allow some kind of abstraction (as in Histogram) to make them reusable * figure out how to do evaluation better (based on the stored abstractions? but need to store provenance...)

SEMANTIC CONDITIONS https://www.youtube.com/watch?v=NBnc2ToS_j0 (has a section on this in background)

INTERACTION * replay stored event handlers against the old version? (this way, adding an item to a speakers list gets migrated & adds a new table row!) * similarly!! we need merge in order to

apply edits to multiple targets (when you remove all items in a list, the indices change) (but I guess we should do this against version at the time of saving too....)

[this & evaluation = the unreasonable effectiveness of merging]

Notes on storing and reusing edits * references need to be represented as references so that they get updated (NO! not if we reply them against old version, which seems better - but there are 2 design choices) * how to apply them to multiple targets? use Move to update the selectors instead of replacing the prefix manually

CONDITIONALS <https://toby.li/files/p311-radensky.pdf>

6.1 Future work

Although Denicek does not explicitly track document structure (or schema, or type), all documents have an implicit structure. => use type system to check temporarily invalid state

better effect system

IDEA: Type check edit groups to ensure they preserve structure but not individual edits eg when adding list item

7 CASE STUDY

(data science environment)

8 EVALUATION

how well did Denicek work for implementing Datnicek

9 DISCUSSION

9.1 Heuristic evaluation

techdms

9.2 Limitations

RANDOM FORMAL STUFF

other formal things * non-conditionality of edits * show that they can transform document from any to any without removing/reading values

PROPERTIES - can change any document to any other - sure, via remove add - but also more semantically (if it contained all values, can we do it without removing and adding them?)

DEMOS

// cf shadow dom v react // something like google forms!

* two ways of pushing edits should be equivalent, i.e.

push e1 e2 e3 over b1 b2 b3 can be done via push e1 over b1 = e1'; over b2 = e1'; over b3 = e1' then push e2, then push e3 or push e1, e2, e3 over b1, then over b2, then over b3 (implies that no "extra" edits can be appended later)

REFERENCES

- [1] Marc T.P. Adam, Shirley Gregor, Alan Hevner, and Stefan Morana. Design science research modes in human-computer interaction projects. *AIS Transactions on Human-Computer Interaction*, 13(1):1–11, 2021.
- [2] Tom Beckmann, Patrick Rein, Stefan Ramson, Joana Bergsieck, and Robert Hirschfeld. Structured editing for all: Deriving usable structured editors from grammars. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, CHI '23. Association for Computing Machinery, 2023.

- [3] A.F. Blackwell. First steps in programming: a rationale for attention investment models. In *Proceedings IEEE 2002 Symposia on Human Centric Computing Languages and Environments*, pages 2–10, 2002.
- [4] Margaret M. Burnett and Brad A. Myers. Future of end-user software engineering: beyond the silos. In *Future of Software Engineering Proceedings*, FOSE 2014, page 201–211. Association for Computing Machinery, 2014.
- [5] Weihao Chen, Xiaoyu Liu, Jiacheng Zhang, Ian Long Lam, Zhicheng Huang, Rui Dong, Xinyu Wang, and Tianyi Zhang. MIWA: mixed-initiative web automation for better user control and confidence. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software Technology*, UIST 2023, San Francisco, CA, USA, pages 75:1–75:15. ACM, 2023.
- [6] James Cheney, Stephen Chong, Nate Foster, Margo Seltzer, and Stijn Vansumeren. Provenance: a future history. In *Proceedings of the 24th ACM SIGPLAN Conference Companion on Object Oriented Programming Systems Languages and Applications*, OOPSLA '09, page 957–964, New York, NY, USA, 2009. Association for Computing Machinery.
- [7] Allen Cypher and Daniel Conrad Halbert. *Watch what I do: programming by demonstration*. MIT press, 1993.
- [8] Andrea A. diSessa and Harold Abelson. Boxer: A reconstructible computational medium. *Commun. ACM*, 29(9):859–868, 1986.
- [9] Ian Drosos, Titus Barik, Philip J. Guo, Robert DeLine, and Sumit Gulwani. Wrex: A unified programming-by-example interaction for synthesizing readable code for data scientists. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, CHI '20, page 1–12. Association for Computing Machinery, 2020.
- [10] Ekwa Duala-Ekoko and Martin P. Robillard. Clonetracker: tool support for code clone management. In *Proceedings of the 30th International Conference on Software Engineering*, ICSE '08, page 843–846, New York, NY, USA, 2008. Association for Computing Machinery.
- [11] Jonathan Edwards. Example centric programming. *SIGPLAN Not.*, 39(12):84–91, December 2004.
- [12] Jonathan Edwards. Subtext: uncovering the simplicity of programming. In *Proceedings of the 20th Annual ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and Applications*, OOPSLA '05, page 505–518, New York, NY, USA, 2005. Association for Computing Machinery.
- [13] Jonathan Edwards. First class copy & paste. Technical Report MIT-CSAIL-TR-2006-037, Massachusetts Institute of Technology, 2006.
- [14] Jonathan Edwards and Tomas Petricek. Interaction vs. abstraction: Managed copy and paste. In *Proceedings of the 1st ACM SIGPLAN International Workshop on Programming Abstractions and Interactive Notations, Tools, and Environments*, pages 11–19, 2022.
- [15] Jonathan Edwards, Tomas Petricek, Tijs van der Storm, and Geoffrey Litt. Schema evolution in interactive programming systems. *The Art, Science, and Engineering of Programming*, 9(?):1–34, 2025.
- [16] Camille Gobert and Michel Beaudouin-Lafon. Lorgnette: Creating malleable code projections. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST '23, New York, NY, USA, 2023. Association for Computing Machinery.
- [17] Sumit Gulwani, William R. Harris, and Rishabh Singh. Spreadsheet data manipulation using examples. *Commun. ACM*, 55(8):97–105, August 2012.
- [18] Christopher Hall, Trevor Standley, and Tobias Hollerer. Infra: structure all the way down: structured data as a visual programming language. In *Proceedings of the 2017 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, Onward! 2017, page 180–197, New York, NY, USA, 2017. Association for Computing Machinery.
- [19] Joshua Horowitz and Jeffrey Heer. Engraff: An api for live, rich, and composable programming. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, UIST '23, New York, NY, USA, 2023. Association for Computing Machinery.
- [20] Joel Jakubovic, Jonathan Edwards, and Tomas Petricek. Technical dimensions of programming systems. *The Art, Science, and Engineering of Programming*, 7(13):1–59, 2023.
- [21] Joel Jakubovic and Tomas Petricek. Ascending the ladder to self-sustainability: Achieving open evolution in an interactive graphical system. In *Proceedings of the 2022 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, Onward! 2022, page 240–258, New York, NY, USA, 2022. Association for Computing Machinery.
- [22] Bas Jansen and Felienne Hermans. Xiblocks: a block-based formula editor for spreadsheet formulas. In *2019 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, pages 55–63, 2019.
- [23] Gowtham Kaki, Swarn Priya, KC Sivaramakrishnan, and Suresh Jagannathan. Mergeable replicated data types. *Proc. ACM Program. Lang.*, 3(OOPSLA), October 2019.
- [24] Sean Kandel, Andreas Paepcke, Joseph Hellerstein, and Jeffrey Heer. Wrangler: interactive visual specification of data transformation scripts. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '11, page 3363–3372, New York, NY, USA, 2011. Association for Computing Machinery.
- [25] Alan C. Kay. The early history of smalltalk. In *The Second ACM SIGPLAN Conference on History of Programming Languages*, HOPL-II, page 69–95. Association for Computing Machinery, 1993.
- [26] Stephen Kell. Unix, plan 9 and the lurking smalltalk. In *Reflections on Programming Systems: Historical and Philosophical Aspects*, pages 189–213. Springer, 2018.
- [27] Stephen Kell and J. Ryan Stinnett. Source-level debugging of compiler-optimised code: Ill-posed, but not impossible. In *Proceedings of the 2024 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, Onward! '24, page 38–53. Association for Computing Machinery, 2024.
- [28] Eugen Kiss. Comparison of object-oriented and functional programming for gui development. Master's thesis, Leibniz Universität Hannover, 2014.
- [29] Cory Kissinger, Margaret Burnett, Simone Stumpf, Neeraja Subrahmanian, Laura Beckwith, Sherry Yang, and Mary Beth Rosson. Supporting end-user debugging: what do users want to know? In *Proceedings of the Working Conference on Advanced Visual Interfaces*, AVI '06, page 135–142. Association for Computing Machinery, 2006.
- [30] Martin Kleppmann, Adam Wiggins, Peter Van Hardenberg, and Mark McGranaghan. Local-first software: you own your data, in spite of the cloud. In *Proceedings of the 2019 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software*, pages 154–178, 2019.
- [31] Clemens Nylandsted Klokmoose, James R Eagan, and Peter van Hardenberg. Mywebstrates: Webstrates as local-first software. In *UIST'24: Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*. ACM, 2024.
- [32] Amy J Ko and Brad A Myers. Designing the whyline: a debugging interface for asking questions about program behavior. In *Proceedings of the SIGCHI conference on Human factors in computing systems*, pages 151–158, 2004.
- [33] Amy J. Ko and Brad A. Myers. Finding causes of program output with the java whyline. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '09, page 1569–1578. Association for Computing Machinery, 2009.
- [34] Eva Krebs, Patrick Rein, Joana Bergsiek, Lina Urban, and Robert Hirschfeld. Probe log: Visualizing the control flow of babylonian programming. In *Companion Proceedings of the 7th International Conference on the Art, Science, and Engineering of Programming*, pages 61–67, 2023.
- [35] Vu Le and Sumit Gulwani. Flashextract: a framework for data extraction by examples. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '14, page 542–553, New York, NY, USA, 2014. Association for Computing Machinery.
- [36] Germán Leiva, Jens Emil Grønbaek, Clemens Nylandsted Klokmoose, Cuong Nguyen, Rubaiat Habib Kazi, and Paul Asente. Rapido: Prototyping interactive experiences through programming by demonstration. In *The 34th Annual ACM Symposium on User Interface Software and Technology*, pages 626–637, 2021.
- [37] Sorin Lerner. Projection boxes: On-the-fly reconfigurable visualization for live programming. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, CHI '20, page 1–7. Association for Computing Machinery, 2020.
- [38] Geoffrey Litt, Sarah Lim, Martin Kleppmann, and Peter van Hardenberg. Peritext: A crdt for collaborative rich text editing. *Proc. ACM Hum.-Comput. Interact.*, 6(CSCW2), November 2022.
- [39] Geoffrey Litt, Peter van Hardenberg, and Henry Orion. Project cambria: Translate your data with lenses. <https://www.inkandswitch.com/cambria.html>, 2020. Accessed: 2020-10-01.
- [40] J. M. Lucassen and D. K. Gifford. Polymorphic effect systems. In *Proceedings of the 15th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL '88, page 47–57. Association for Computing Machinery, 1988.
- [41] Sean McDirmid. Usable live programming. In *Proceedings of the 2013 ACM International Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software*, Onward! 2013, page 53–62. Association for Computing Machinery, 2013.
- [42] Robert C. Miller and Brad A. Myers. Interactive simultaneous editing of multiple text regions. In *Proceedings of the General Track: 2001 USENIX Annual Technical Conference*, page 161–174, USA, 2001. USENIX Association.
- [43] Brad A. Myers, Amy J. Ko, and Margaret M. Burnett. Invited research overview: end-user programming. In *CHI '06 Extended Abstracts on Human Factors in Computing Systems*, CHI EA '06, page 75–80. Association for Computing Machinery, 2006.
- [44] Brad A. Myers, Richard McDaniel, and David Wolber. Programming by example: intelligence in demonstrational interfaces. *Commun. ACM*, 43(3):82–89, March 2000.
- [45] Bonnie A. Nardi and James R. Miller. The spreadsheet interface: A basis for end user programming. In *Proceedings of the IFIP TC13 Third International Conference on Human-Computer Interaction*, INTERACT '90, page 977–983, NLD, 1990. North-Holland Publishing Co.
- [46] Cyrus Omar, David Moon, Andrew Blinn, Ian Voysey, Nick Collins, and Ravi Chugh. Filling typed holes with live guis. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, PLDI 2021, page 511–525, New York, NY, USA, 2021. Association for Computing

- Machinery.
- [47] Cyrus Omar, Ian Voysey, Michael Hilton, Jonathan Aldrich, and Matthew A. Hammer. Hazelnut: a bidirectionally typed structure editor calculus. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*, POPL '17, page 86–99, New York, NY, USA, 2017. Association for Computing Machinery.
 - [48] Addy Osmani, Sindre Sorhus, Pascal Hartig, and Stephen Sawchuk. Todomvc: Helping you select an MV* framework. <https://todomvc.com/>, 2024. Accessed: 2024-12-12.
 - [49] Roly Perera, Umut A. Acar, James Cheney, and Paul Blain Levy. Functional programs that explain their work. In *Proceedings of the 17th ACM SIGPLAN international conference on Functional programming*, pages 365–376, 2012.
 - [50] Roly Perera, Minh Nguyen, Tomas Petricek, and Meng Wang. Linked visualisations via galois dependencies. *Proceedings of the ACM on Programming Languages*, 6(POPL):1–29, 2022.
 - [51] Tomas Petricek. Foundations of a live data exploration environment. *The Art, Science, and Engineering of Programming*, 4(8):1–37, 2020.
 - [52] Tomas Petricek. Foundations of a live data exploration environment. *The Art, Science, and Engineering of Programming*, 4(8):1–37, 2020.
 - [53] Erhard Rahm and Philip A. Bernstein. An online bibliography on schema evolution. *SIGMOD Rec.*, 35(4):30–31, December 2006.
 - [54] David Rauch, Patrick Rein, Stefan Ramson, Jens Lincke, and Robert Hirschfeld. Babylonian-style programming: Design and implementation of a general-purpose editor integrating live examples into source code. *The Art, Science, and Engineering of Programming*, 3(9):1–39, 2019.
 - [55] Patrick Rein, Stefan Ramson, Jens Lincke, Robert Hirschfeld, and Tobias Pape. Exploratory and live, programming and coding: A literature study comparing perspectives on liveness. *The Art, Science, and Engineering of Programming*, 3(1):1–33, 2019.
 - [56] Wilmer Ricciotti, Jan Stolarek, Roly Perera, and James Cheney. Imperative functional programs that explain their work. *Proceedings of the ACM on Programming Languages*, 1(ICFP):1–28, 2017.
 - [57] Advait Sarkar, Andrew D. Gordon, Simon Peyton Jones, and Neil Toronto. Calculation view: multiple-representation editing in spreadsheets. In *2018 IEEE Symposium on Visual Languages and Human-Centric Computing (VL/HCC)*, pages 85–93, 2018.
 - [58] David Canfield Smith. *Pygmalion: A Creative Programming Environment*. PhD thesis, Stanford University, 1975. Available as Stanford AI Memo AIM-260 and Computer Science Report STAN-CS-75-499.
 - [59] Randall B. Smith, John Maloney, and David Ungar. The self-4.0 user interface: manifesting a system-wide vision of concreteness, uniformity, and flexibility. In *Proceedings of the Tenth Annual Conference on Object-Oriented Programming Systems, Languages, and Applications*, OOPSLA '95, page 47–60, New York, NY, USA, 1995. Association for Computing Machinery.
 - [60] Guy L. Steele and Richard P. Gabriel. The evolution of lisp. In *The Second ACM SIGPLAN Conference on History of Programming Languages*, HOPL-II, page 231–270. Association for Computing Machinery, 1993.
 - [61] M. Toomim, A. Begel, and S.L. Graham. Managing duplicated code with linked editing. In *2004 IEEE Symposium on Visual Languages - Human Centric Computing*, pages 173–180, 2004.

A FORMATIVE EXAMPLES

The design the Denicek substrate, we identified six formative examples shown in Fig. 15. The examples range from established industry benchmarks (Todo and Counter apps) to cases from literature [14] and problems posed as schema change challenges [15]. The Denicek substrate then co-evolved with Webnicek, a simple web-based programming environment built (as directly as possible) on top of the substrate and was used to solve implement the formative examples.

Many of the formative examples include a small programming challenge, such as adding user interface to add a new speaker, a new list item or modify the count. Our aim was for the substrate to enable solving those through programming by demonstration. Programming by Demonstration is often used in data wrangling [9, 17, 35]. Our Hello World example is only a minimalistic illustration of such use, loosely inspired by earlier work [42].

B MERGING EDIT HISTORIES

Recall that merging takes two edit histories, E, E_1 and E, E_2 , transforms edits E_2 into E'_2 that can be reapplied on top of the first history resulting in E, E_1, E'_2 . The key operation takes two individual edits, e_1 and e_2 and produces a sequence of edits e'_2, e'_2, \dots that can be applied after e_1 , and combine the two edits. This section provides details about the two aspects of this operation.

B.1 Apply to Newly Added

Assume that edits e_1 and e_2 occurred independently. We want to modify e_2 so that it can be placed after e_1 . If the edit e_2 added new nodes to the document that the edit e_1 would affect, we generate an additional edit that apply the transformation of e_1 to the newly added nodes (and only to those).

The only edits that add new document nodes are Add, Append, Copy and so we consider this case if the edit e_2 is one of those. If so, we check whether the target of e_1 is within the target of e_2 , i.e., the list of selectors that forms the target of e_2 is a prefix of the list of selectors that forms the target of e_1 .

Along the way, we compute a *more specific prefix*. If the target of e_1 contains the All selector, it can be matched against a specific Index selector in the target of e_2 (if the selector of e_1 is more specific than that of e_2 , the targets are not matched). We then replace the original prefix in e_1 with the *more specific prefix* that contains Index selector in places where the original edit contained All. This way, we obtain e'_1 which is a focused version of e_1 that applies only to the nodes newly added by e_2 . The edit e_2 thus becomes a pair of edits e_2, e'_1 . The final document will contain edits e_1, e_2, e'_1 – that is, it will first apply the edit e_1 to nodes already in the document, then add new nodes and then apply the transformation represented by e_1 to the newly added nodes.

B.2 Transform Matching References

As above, assume that edits e_1 and e_2 occurred independently. We want to modify e_2 so that it can be placed after e_1 . If e_1 is any of the three edits listed in Fig. 4 (RenameField, WrapRecord, WrapList), we collect all references that appear inside e_2 (the target, the source of Copy and any references occurring in the nodes added by Add or Append). If the target of e_1 is a prefix of any of those references, we update the references accordingly and obtain a new edit e'_2 . Note

Counter App [28] – Counter with increment and decrement buttons.

The current count is represented by a formula that is modified by the buttons. The user can inspect the evaluation trace to see how the count was modified.

Programming by Demonstration, Incremental Recomputation, End-User Debugging

Todo App [48] – Buttons to add an item and remove all completed.

Adding an item must correctly merge with independently added functionality to compute which items are completed and remove them based on a formula result.

Programming by Demonstration, Local-First Collaboration, Incremental Recomputation

Conference List [15] – Manage a list of invited conference speakers.

Adding speakers to a list through an in-document user interface merges with refactoring that turns the list into a table and separates name from an email.

Local-First Collaboration, Programming by Demonstration

Conference Budget [15] – Calculate budget based on a speaker list.

References are updated when the list is refactored. Only affected formulas are recomputed and the user can view elements on which the result depends.

Local-First Collaboration, Incremental Recomputation, End-User Debugging

Hello World [42] – Normalize the capitalization of two word messages.

An operation to normalize the text in a list item can be recorded and applied to all list items or, alternatively, applied directly to all list items.

Programming by Demonstration

Traffic Accidents [14] – Compute statistics using two data sources.

Formula to compute statistics can be reused with a different data source; error correction is propagated automatically to the copied version of the formula.

Concrete Programming, Incremental Recomputation

Figure 15: Formative examples used in Denicek design

that it would be an error to match specific Index in e_1 with more general All in e_2 , but this cannot happen – reference updating is not done when the target of e_1 contains Index.

Now consider the case when e_1 is Copy and the edit e_2 targets a node that is the source node of the copy operation (or any of its children). In this case, it is reasonable to require that the edit e_2 is applied to both the source and the target of the copy. (This is required by the refactoring done in the Conference Budget example.) We handle the case by creating a copy of e_2 with transformed selectors (target and, if e_2 is also Copy, also its source). To transform the selectors, we replace the prefix formed by the source of the Copy by a new prefix, formed by the target target of the Copy. We then add the new operation as e'_2 if at least one of its selectors was transformed (typically target, but possibly also source).

XX

introduction
 background
 related work
 design process / goals
 case study
 formative research
 formative study
 analysis
 system
 implementation
 evaluation / heuristic evaluation
 discussion and limitations

C INTRODUCTION OLD

The computational substrate using which software is built determines the capabilities that the software can provide. An imperative substrate that views programs as instructions modifying bytes in memory makes it almost impossible to allow end-user inspection or reprogramming of running software.

substrate as defined by [21]

A computational substrate defines what software is built from. This may be objects as in Smalltalk, lists as in Lisp, or memory with data and code as in UNIX/C. The different substrates enable different kinds of programming experiences. For example, object-oriented programming has historically been linked to the development of graphical user interfaces (where objects can correspond to elements on the screen). It has also enabled the development of visual programming environments such as the Alternate Reality Kit, based on message sending between objects.

In principle, any computational substrate can be used to develop any programming experience, but the greater the impedance mismatch between the substrate and the desired experience, the more difficult it will be to provide the experience and combine it with the rest of the system and other programming experiences developed for the system. (One can implement support for programming-by-demonstration using C/C++, for example as part of a game scripting engine, but it will not work with the rest of the ordinary C/C++ ecosystem.)

Medium is the message

C.1 Substrate

The question asked in this paper is, what would be the ideal programming substrate for supporting a range of programming experiences that make programs more collaborative, transparent and allows for a gradual transition from non-programmer to a programmer. We want a programming substrate that makes it easy to develop programming experiences such as:

- *Programming by demonstration* – Allow non-programmers to construct simple programs by performing examples of the expected behaviour. [36].
- *Local-first collaboration* – Multiple users should be able to use and modify a single program, preferably without requiring a central server. [30]
- *Provenance tracking* – The execution of the program should leave an understandable trace that lets the user understand why program resulted in a particular result.
- *Schema evolution [extra-ish]* – When the user evolves the structure of the program, data and code should co-evolve automatically to match the new structure.
- *Notational freedom [extra-ish?]* – Allow users to adapt the program using a notation that suits them and is appropriate for the programming task at hand. [Joel]
- *Concrete programming [extra?]* – It should be possible to reuse parts of program or program logic without constructing abstractions, for example by managed copy & paste.[13, 14]

substrate as defined by [21]

[7, 36] [5]

[30, 31] [32–34] [49, 56] [50]

Joel's definition of substrate in Onward! Bret Victor talk <https://www.youtube.com/watch?v=ef2jpiTEB5U>

In what ways is a substrate "natural"?

thinglab - create line by cloning, it sticks to mouse pointer, clicking sticks it to something else squeak - has all the browsers (method search...)

computational substrate how it differs from computational media? more low-level - media suggests that there it comes

D THE WHATEVER SYSTEM

D.1 Document + Edits

defines

- selectors
- nodes
- edits

D.2 Walkthrough

* todo list? (or counter, but that is a bit boring)

E THEMES

* programming by demonstration - binding interactions to gui elements (event handlers) * provenance tracking - Amy Ko's whyline, Probe Log by HPI, enables linked visualizations * merging of edit histories / collaborative editing - bonus - can share restricted link to allow users fill out forms (allow partial edits only / def by selector?) * schema change - change data & code accordingly * everything is an edit - interaction with the GUI - evaluation? tbd * copy & paste abstraction (requires finishing new approach to formulas!) - edit before copy to propagate edit to other places (or edit after copy to make it specific to a case) - higher order copying from <https://tomasp.net/academic/papers/copy-paste/paint22.pdf> * augmenters - cf. bonnie Nardi (calls them something else - Jonathan says) [data detectors https://dl.acm.org/doi/10.1145/272287.272331](https://dl.acm.org/doi/10.1145/272287.272331) also see Potluck <https://www.inkandswitch.com/potluck/> - add programming by demonstration data wrangling gui to table (trigger interactions) cf. lorgnette

F APPLICATIONS

* todo list / counter / maybe too simple * (if used in the walkthrough, maybe something else? board game as in varv - tic tac toe? or 7guis?) * conference organizer * data exploration (ala histogram) * linked charts

G EXTRAS

* metablocks? * self-sustainability * some non-browser implementation of this (as in Varv?)

explicit structure self-sustainability notational freedom

REMAINING IMPLEMENTATION TODOs:

* Some kind of provenance visualization * Some kind of matchers/transformers mechanism (ideally to add interactive buttons to tables) * Apply to all (remove completed in TODO)


```

append /speakers #hamilton <li> addfld /speakers/#hamilton
speaker "Margaret Hamilton, hamilton@mit.edu"
  »> wraprec /speakers/body <table>
    [update selectors - add /body] append /speakers/body #hamil-
ton <li> [ditto] addfld /speakers/body#hamilton speaker "Margaret
Hamilton, hamilton@mit.edu"
  »> retag /speakers/body <tbody>
    [na - out of scope]
  »> retag /speakers/body/* <tr>
    [apply to added - update node (* => #hamilton scoping)] append
/speakers/body #hamilton <li> + retag /speakers/body/#hamilton
<tr> [na - out of scope] addfld /speakers/body#hamilton speaker
"Margaret Hamilton, hamilton@mit.edu"
  »> wraprec /speakers/body/* /speaker value <td>
    append /speakers/body #hamilton <li> + retag /speakers/body/#hamilton
<tr> + wraprec /speakers/body/#hamilton/speaker value <td> [ no
op - it is not there] addfld /speakers/body/#hamilton speaker "Mar-
garet Hamilton, hamilton@mit.edu" + wraprec /speakers/body/#hamilton/
value <td>

```

```

  »> updid /speakers/body/* speaker=>name
    append /speakers/body #hamilton <li> + retag /speakers/body/#hamilton
<tr> + wraprec /speakers/body/#hamilton/speaker value <td> [ no
op - it is not there] + updid /speakers/body/#hamilton speaker=>name
[ no op - it is not there ] addfld /speakers/body #hamilton speaker
"Margaret Hamilton, hamilton@mit.edu" + wraprec /speakers/body/#hamilton/spea-
value <td> + updid /speakers/body/#hamilton speaker=>name
  »> addfld /speakers/body/* email ""
    append /speakers/body #hamilton <li> + retag /speakers/body/#hamilton
<tr> + wraprec /speakers/body/#hamilton/speaker value <td> [ no
op - it is not there] + updid /speakers/body/#hamilton speaker=>name
[ no op - it is not there ] + addfld /speakers/body/#hamilton email ""
[ no op - it is not there ] addfld /speakers/body #hamilton speaker
"Margaret Hamilton, hamilton@mit.edu" + wraprec /speakers/body/#hamilton/spea-
value <td> + updid /speakers/body/#hamilton speaker=>name +
    addfld /speakers/body/#hamilton email ""
  ...
    copy /speakers/body/* /name => /speakers/body/* /email edit /speak-
ers/body/* /name (fn1) edit /speakers/body/* /email (fn2)

```